

```
In [15]: import pandas as pd  
from matplotlib import pyplot as plt  
import seaborn as sns  
import statsmodels.formula.api as smf
```

## 1. Importing Data

```
In [16]: salary_data=pd.read_csv('Salary_Data.csv')
salary_data
```

```
Out[16]:
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	56642.0
6	3.0	60150.0
7	3.2	54445.0
8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55794.0
12	4.0	56957.0
13	4.1	57081.0
14	4.5	61111.0
15	4.9	67938.0
16	5.1	66029.0
17	5.3	83088.0
18	5.9	81363.0
19	6.0	93940.0
20	6.8	91738.0
21	7.1	98273.0
22	7.9	101302.0
23	8.2	113812.0
24	8.7	109431.0
25	9.0	105582.0
26	9.5	116969.0
27	9.6	112635.0
28	10.3	122391.0
29	10.5	121872.0

```
In [17]: salary_data.shape
```

```
Out[17]: (30, 2)
```

```
In [18]: salary_data.head()
```

```
Out[18]:
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0

```
In [19]: salary_data.isna().sum()
```

```
Out[19]: YearsExperience    0
Salary                    0
dtype: int64
```

```
In [20]: salary_data.dtypes
```

```
Out[20]: YearsExperience    float64
Salary                    float64
dtype: object
```

```
In [21]: salary_data.describe()
```


```
Out[21]:
```

	YearsExperience	Salary
count	30.000000	30.000000
mean	5.313333	76003.000000
std	2.837888	27414.429785
min	1.100000	37731.000000
25%	3.200000	56720.750000
50%	4.700000	65237.000000
75%	7.700000	100544.750000
max	10.500000	122391.000000

## 2. Check for assumptions

### 1. linearity check

```
In [22]: sns.scatterplot(x='YearsExperience',y='Salary',data=salary_data)  
plt.show()
```




```
In [23]: salary_data.corr()
```

```
Out[23]:
```

	YearsExperience	Salary
YearsExperience	1.000000	0.978242
Salary	0.978242	1.000000

```
In [24]: sns.regplot(x='YearsExperience',y='Salary',data=salary_data)  
plt.show()
```



### 3.Model Building || Model Training

```
In [25]: #model training
lin_model=smf.ols(formula='Salary~YearsExperience',data=salary_data).fit()
lin_model
```

```
Out[25]: <statsmodels.regression.linear_model.RegressionResultsWrapper at 0x17586620d30>
```

## 4.Model Testing

```
In [26]: lin_model.params
```

```
Out[26]: Intercept          25792.200199
YearsExperience      9449.962321
dtype: float64
```

```
In [27]: lin_model.tvalues,lin_model.pvalues
```

```
Out[27]: (Intercept          11.346940
YearsExperience      24.950094
dtype: float64,
Intercept          5.511950e-12
YearsExperience      1.143068e-20
dtype: float64)
```

```
In [28]: lin_model.rsquared
```

```
Out[28]: 0.9569566641435086
```

```
In [29]: # Machine prediction
pred_data={'YearsExperience':[2,5,7]}
pred_data
```

```
Out[29]: {'YearsExperience': [2, 5, 7]}
```

```
In [30]: test_data=pd.DataFrame(data=pred_data)
test_data
```

```
Out[30]:
```

	YearsExperience
0	2
1	5
2	7

```
In [31]: lin_model.predict(test_data)
```

```
Out[31]: 0    44692.124842
1    73042.011806
2    91941.936449
dtype: float64
```

